EMERGE AI HACKATHON · HEALTHCARE · CHALLENGE 5 · SAT OCT 17 2026

Hear Me Out

Never miss the sentence a person needs to see.

A second listener for health-coaching calls (eMerge Challenge 5, Coach's Copilot). Reads bilingual health-coaching calls and flags the sentence a person needs to see today.

Esteban Guio, solo builder, with an AI team

THE PROBLEM

A coach hears it. The call moves on.

A coach is on a call with a cancer survivor. In Spanish, the person says “No puedo respirar bien.” The coach, with a full caseload, moves on to meal planning. Nobody listens to that call for weeks. **Hear Me Out is a second listener.**

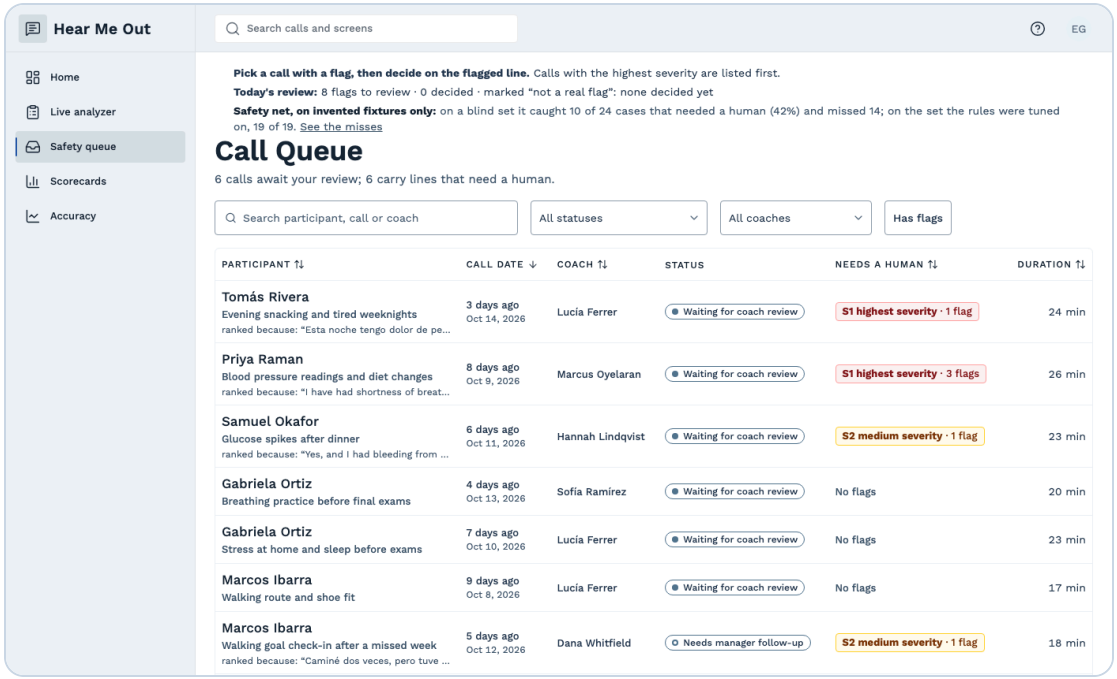
Illustrative scenario, not a real patient.

- Coaching calls are hard to review at scale, so safety signals can slip by.
- Coaches vary in how they coach, and supervisors have little visibility.
- Spanish-language calls are the hardest to review at scale. We found no Spanish-first coaching review tool in our search today.

Problem framing from the challenge brief; no statistic is claimed on this slide.

WHAT IT DOES

Tag the call. Score the coach. Flag the human moment.



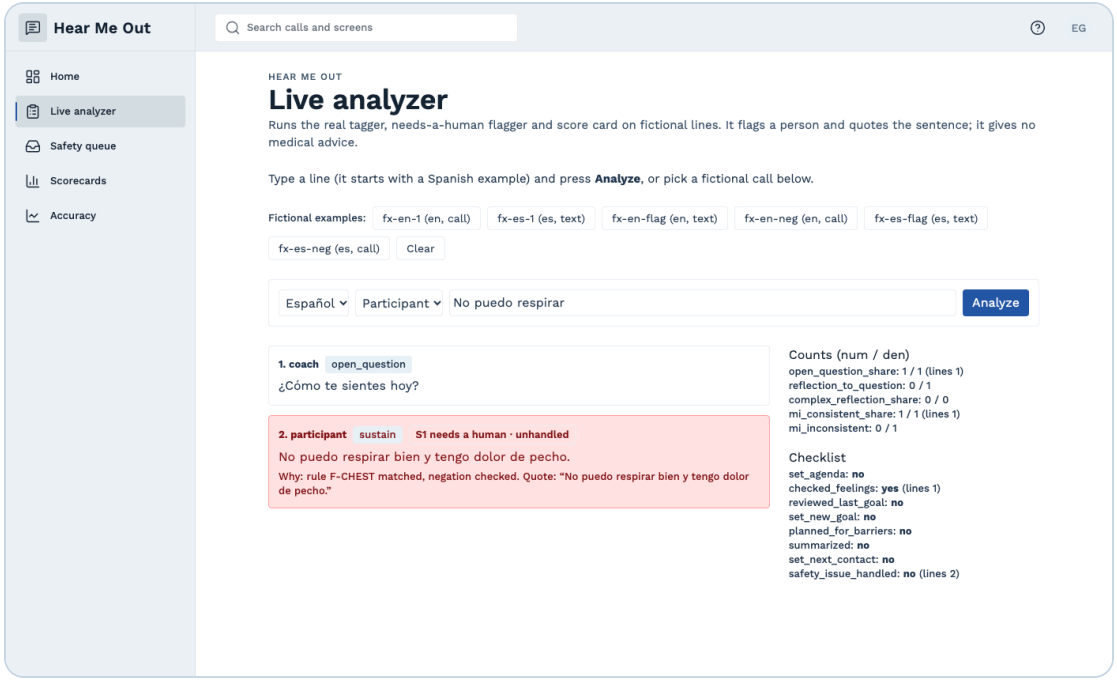
- **Safety queue:** calls that need a human today, highest severity first.
- **Quoted sentence:** each flag shows the exact line and why.
- **Coach score card:** motivational-interviewing counts with the evidence lines.
- No diagnosis. No medical advice.

Screenshot: live app at coachs-copilot.onrender.com, Oct 17 2026. All people and calls shown are invented.

LIVE DEMO

Spanish in, flagged sentence out.

Think of a smoke detector, not a firefighter: it doesn't diagnose or put anything out; it makes sure a human looks, today.



- A fictional Spanish call: **"No puedo respirar bien y tengo dolor de pecho."**
- The real engine flags it: **S1, needs a human, unhandled**, with the rule and the quoted sentence.
- Then a denial: "no tengo fiebre" stays quiet; "no puedo respirar" does not. Negation is handled.

Fictional lines only. No challenge data is in the app or the public repo.

What we used, why, and how we checked it.

TOOL	BUILT	WHY THIS TOOL	HOW WE CHECKED IT
Claude Code	Engine (loader, tagger, flagger, score card), API, tests, deploy prep	One owner of main; tests with every module	113 engine tests and 18 eval tests pass
Codex	Read-only code reviews, several rounds	An independent second reader that cannot edit	High-severity findings fixed before merge
Devin	Offline evaluation harness (precision, recall, kappa)	Isolated chunk, mock data only	Pull request merged after review
Archie	The interface, from a data-free contract	Fastest route to a usable UI	Wired to our engine, tested in the browser
Render	Free hosting of the public demo	One URL judges can open	Invented data only

Rules, not a model: the safety path never depends on an AI call.

A stress test of about 11,000 requests caused no crashes and found three gaps (extra whitespace, full-width text, Spanish crisis phrases). They were then patched.

Source: the repository's AI_TOOLS.md and reviews folder; test counts from a run on Oct 17 2026, 14:2x.

THE HONEST NUMBERS

Two results. The second one is the honest one.

TEST SET (INVENTED TEXT)	CAUGHT (S1–S3)	CAUGHT (S1 ONLY)	FALSE ALARMS PER 10 CALLS
Tuned set: 28 cases, rules adjusted while looking	100%	100%	0
Blind set: 36 new phrasings, run once	**42%** (10 of 24)	**50%** (5 of 10)	0.56

English caught 3 of 12. Spanish caught 7 of 12. Small samples.

The rules catch common wording and miss indirect, informal phrasing. That is why a human decides every flag.

Source: the app's Accuracy page, read Oct 17 2026. Invented text only; not a claim about real conversations. Blind figures were measured on the rules as frozen before the later patches; the patches are tested separately.

Limits

- A safety net, not a diagnosis.
- Every flag is reviewed by a human.
- Invented data only so far.
- Spanish and dialect coverage is the main known gap.
- MI scores are a coaching aid, not for HR action.
- Decisions are logged.

Built

- Export of each decided flag as a FHIR Flag file (demo structure, not validated against an EHR profile).
- A decision-log download.
- A local mode with no outside calls.

WHAT IT TAKES TO RUN THIS FOR REAL

Three steps. No promises of results.

1. Run it where the data lives
The engine has a local mode with no outside calls, so a coaching program runs it on its own machine and nothing leaves. The public demo uses invented fixtures only.

2. Clinical review before real use
A nurse and a native Spanish speaker review the flag rules. Today they are our own heuristics, the S1 to S3 tiers are our design, and the blind figure is measured on invented phrasings.

3. Next build
Per-line language detection for mixed English and Spanish calls, Kreyòl patterns, a sign-in and a decision log (the FHIR Flag export and decision-log download are the first step), then a validation run on real, consented data, reported as real numbers.

Human decides. It never diagnoses.

ASK AND SOURCES

We want one pilot partner.

A coaching or survivorship program willing to review a handful of real calls with us, under their own privacy rules.

Built from the challenge brief and dataset documentation. Not yet validated with real users or in a pilot.

Sources and credits. Challenge brief and dataset documentation: eMerge AI Hackathon, Miami AI Hub, Sylvester Comprehensive Cancer Center. Dataset used in our local evaluation (not published): AnnoMI: Wu et al., “Anno-MI: A Dataset of Expert-Annotated Counselling Dialogues,” ICASSP 2022, pp. 6177–6181. The public repository contains code only, no challenge data.

Try it: https://coachs-copilot.onrender.com · **Code:** https://github.com/theguiotine/hackathon-copilot